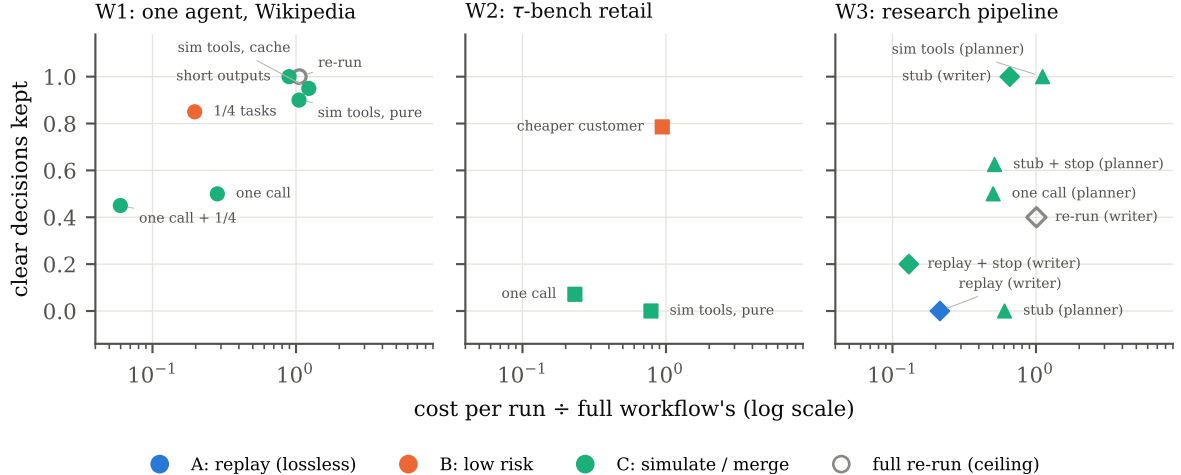
W1: one agent, Wikipedia
W2: τ-bench retail
W3: research pipeline
clear decisions kept
1.0
0.8
0.6
0.4
0.2
0.0
sim tools, cache
re-run
short outputs
1/4 tasks
sim tools, pure
one call
one call + 1/4
cheaper customer
one call
sim tools, pure
sim tools (planner)
stub (writer)
stub + stop (planner)
one call (planner)
re-run (writer)
replay + stop (writer)
replay (writer)
stub (planner)
10^{-1}
10^{0}
cost per run ÷ full workflow's (log scale)
A: replay (lossless)
B: low risk
C: simulate / merge
full re-run (ceiling)