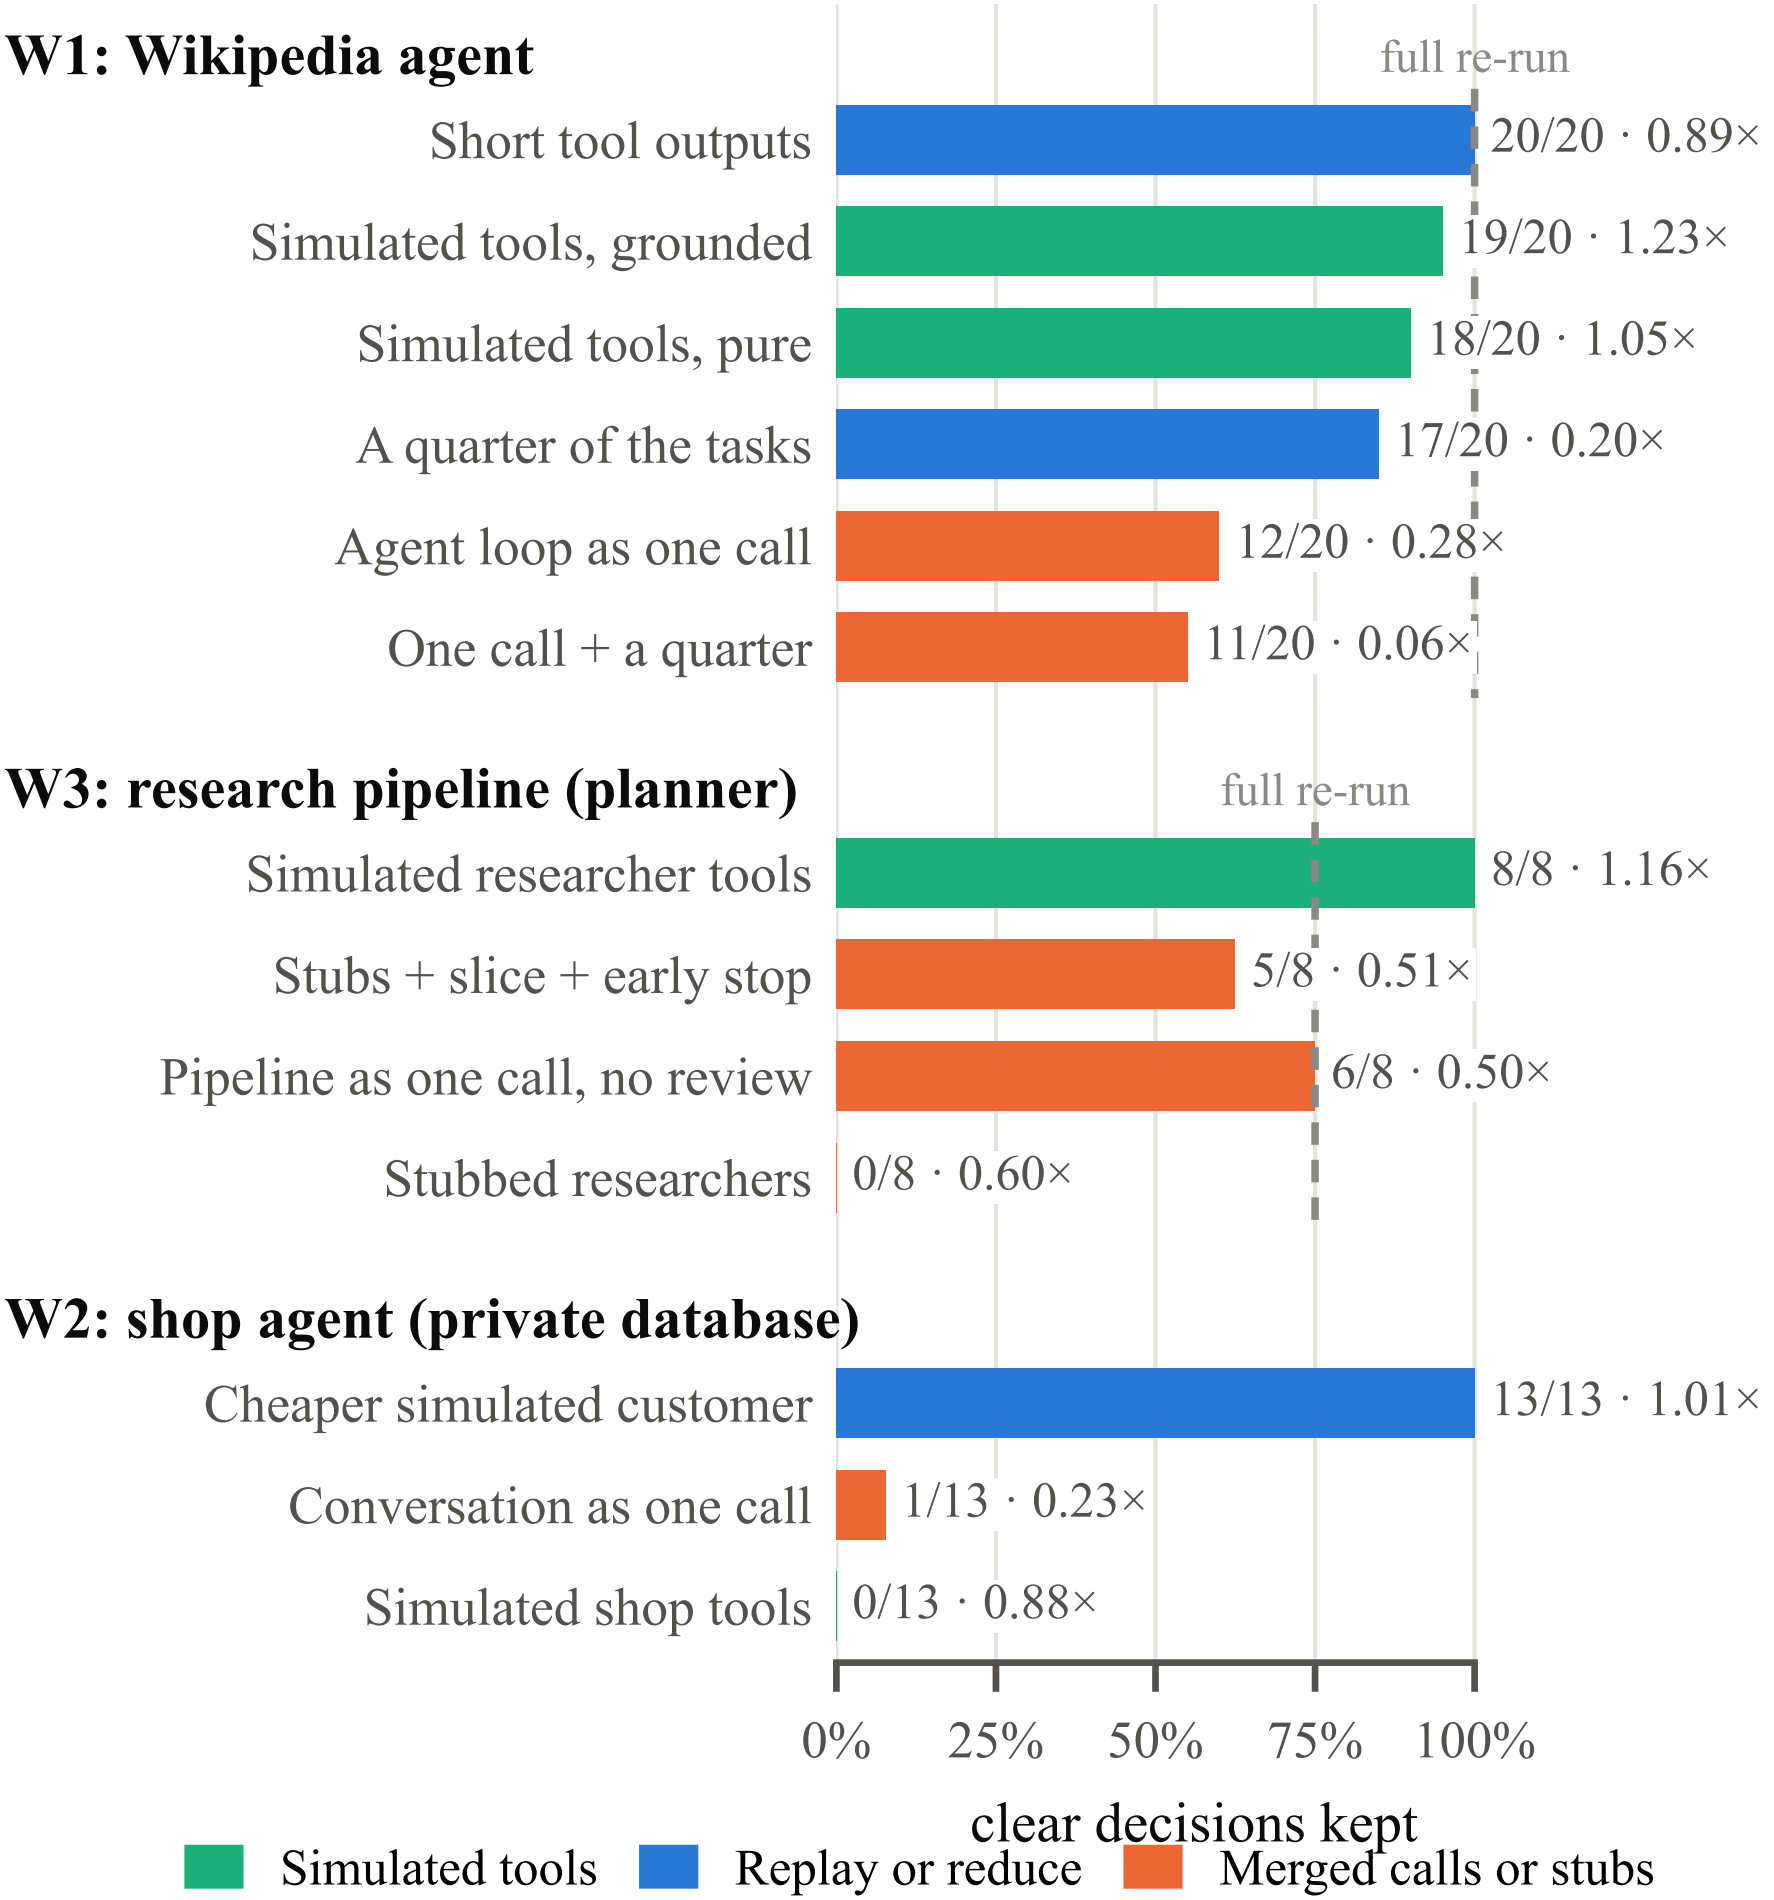
W1: Wikipedia agent
full re-run
Short tool outputs
20/20 · 0.89×
Simulated tools, grounded
19/20 · 1.23×
Simulated tools, pure
18/20 · 1.05×
A quarter of the tasks
17/20 · 0.20×
Agent loop as one call
12/20 · 0.28×
One call + a quarter
11/20 · 0.06×
W3: research pipeline (planner)
full re-run
Simulated researcher tools
8/8 · 1.16×
Stubs + slice + early stop
5/8 · 0.51×
Pipeline as one call, no review
6/8 · 0.50×
Stubbed researchers
0/8 · 0.60×
W2: shop agent (private database)
Cheaper simulated customer
13/13 · 1.01×
Conversation as one call
1/13 · 0.23×
Simulated shop tools
0/13 · 0.88×
0%
25%
50%
75%
100%
clear decisions kept
Simulated tools
Replay or reduce
Merged calls or stubs